






# HW5

Keeno Glanville

# Data Exploration

The data exploration was quite telling, it was able to alleviate personal expectations about age distribution. I would assume more youth would be involved in the crashes. The data was also not strongly correlated which would be intuitive seeing as no one can really predict a crash



# Data Preparation

TARGET_FLAG	0	TARGET_FLAG	0
TARGET_AMT	0	TARGET_AMT	0
KIDSDRIV	0	KIDSDRIV	0
AGE	6	AGE	0
HOMEKIDS	0	HOMEKIDS	0
YOJ	454	YOJ	0
INCOME	445	INCOME	0
PARENT1	0	PARENT1	0
HOME_VAL	464	HOME_VAL	0
MSTATUS	0	MSTATUS	0
SEX	0	SEX	0
EDUCATION	0	EDUCATION	0
JOB	526	JOB	0
TRAVTIME	0	TRAVTIME	0
CAR_USE	0	CAR_USE	0
BLUEBOOK	0	BLUEBOOK	0
TIF	0	TIF	0
CAR_TYPE	0	CAR_TYPE	0
RED_CAR	0	RED_CAR	0
OLDCLAIM	0	OLDCLAIM	0
CLM_FREQ	0	CLM_FREQ	0
REVOKED	0	REVOKED	0
MVR_PTS	0	MVR_PTS	0
CAR_AGE	510	CAR_AGE	0
URBANICITY	0	URBANICITY	0
dtype: int64		dtype: int64	



The data preparation consisted of fixing typographical inconsistencies as well as imputing missing values. For the most part the dataset was easy to work with.

```
# Convert object columns to numeric
train['INCOME'] = pd.to_numeric(train['INCOME'].str.replace('[\$', ' '), errors='coerce')
train['HOME_VAL'] = pd.to_numeric(train['HOME_VAL'].str.replace('[\$', ' '), errors='coerce')
train['BLUEBOOK'] = pd.to_numeric(train['BLUEBOOK'].str.replace('[\$', ' '), errors='coerce')
train['OLDCLAIM'] = pd.to_numeric(train['OLDCLAIM'].str.replace('[\$', ' '), errors='coerce')
train['MSTATUS'] = train['MSTATUS'].replace('z_No', 'No')
train['SEX'] = train['SEX'].replace('z_F', 'F')

# Convert object columns to numeric
test['INCOME'] = pd.to_numeric(test['INCOME'].str.replace('[\$', ' '), errors='coerce')
test['HOME_VAL'] = pd.to_numeric(test['HOME_VAL'].str.replace('[\$', ' '), errors='coerce')
test['BLUEBOOK'] = pd.to_numeric(test['BLUEBOOK'].str.replace('[\$', ' '), errors='coerce')
test['OLDCLAIM'] = pd.to_numeric(test['OLDCLAIM'].str.replace('[\$', ' '), errors='coerce')
test['MSTATUS'] = test['MSTATUS'].replace('z_No', 'No')
test['SEX'] = test['SEX'].replace('z_F', 'F')

# Print updated column types
column_types = train.dtypes
print(column_types)
```

# Build Models

```
# Define features and target
X_flag = train.drop(['TARGET_FLAG', 'TARGET_AMT'], axis=1)
y_flag = train['TARGET_FLAG']

# Convert categorical variables to dummies
X_flag = pd.get_dummies(X_flag, drop_first=True)

# Split the data into training and testing sets
X_train_flag, X_test_flag, y_train_flag, y_test_flag = train_test_split(X_flag, y_flag, test_size=0.2, random_state=42)

# Binary Logistic Regression
log_reg_flag = LogisticRegression(random_state=42)
log_reg_flag.fit(X_train_flag, y_train_flag)

# Evaluate the model
y_pred_flag = log_reg_flag.predict(X_test_flag)
print(confusion_matrix(y_test_flag, y_pred_flag))
print(classification_report(y_test_flag, y_pred_flag))
```

```
[[1158  31]
 [ 418  26]]

      precision    recall  f1-score   support

     0       0.73       0.97       0.84       1189
     1       0.46       0.06       0.10        444

 accuracy          0.73       1633
 macro avg       0.60       0.52       0.47       1633
 weighted avg    0.66       0.73       0.64       1633
```

```
# Define features and target
X_amt = train.drop(['TARGET_FLAG', 'TARGET_AMT'], axis=1)
y_amt = train['TARGET_AMT']

# Convert categorical variables to dummies
X_amt = pd.get_dummies(X_amt, drop_first=True)

# Split the data into training and testing sets
X_train_amt, X_test_amt, y_train_amt, y_test_amt = train_test_split(X_amt, y_amt, test_size=0.2, random_state=42)

# Linear Regression for predicting continuous variable
lin_reg_amt = LinearRegression()
lin_reg_amt.fit(X_train_amt, y_train_amt)

# Evaluate the model
y_pred_amt = lin_reg_amt.predict(X_test_amt)
print('Mean Squared Error:', mean_squared_error(y_test_amt, y_pred_amt))
print('R-squared:', r2_score(y_test_amt, y_pred_amt))
```

Mean Squared Error: 28992700.242169943

R-squared: 0.061593652624510664

# Select Models

Model was tuned to select the most optimal features to present a more accurate model, however the model had a very difficult time fitting the data

```
# Linear Regression with Feature Selection
lin_reg_amt = LinearRegression()

# Use SelectFromModel for feature selection
feature_selector_amt = SelectFromModel(lin_reg_amt)
X_train_amt_selected = feature_selector_amt.fit_transform(X_train_amt, y_train_amt)
X_test_amt_selected = feature_selector_amt.transform(X_test_amt)

# Train the model with selected features
lin_reg_amt.fit(X_train_amt_selected, y_train_amt)

# Evaluate the model
y_pred_amt = lin_reg_amt.predict(X_test_amt_selected)
print('Mean Squared Error:', mean_squared_error(y_test_amt, y_pred_amt))
print('R-squared:', r2_score(y_test_amt, y_pred_amt))
```

Mean Squared Error: 29467483.956893705

R-squared: 0.04622633437523449